

Introduction to R: Computation & Regression Analysis

Kory Kantenga

University of Pennsylvania

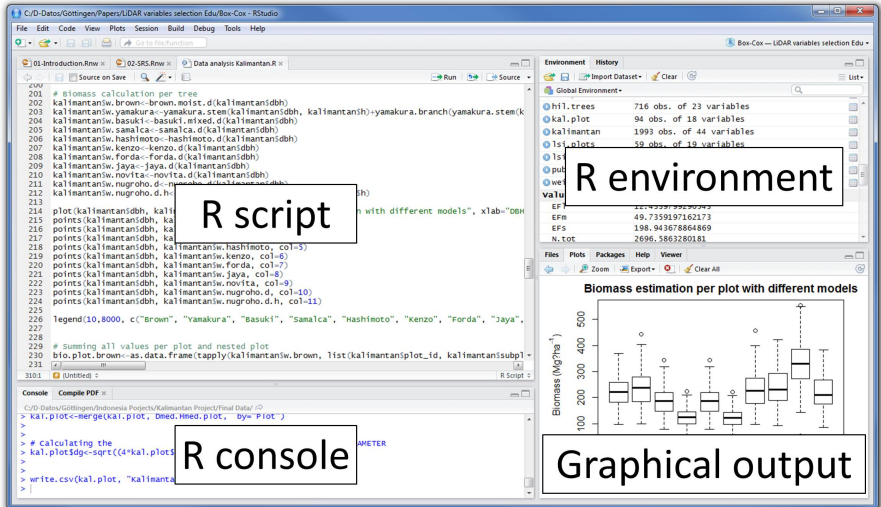
1. What is R?
2. Computing in R
3. Graphics in R
4. More in R

R is a language for programming (e.g. statistical, graphical)



R Studio is software application in which we can execute *R* code





The screenshot displays the R Studio environment with three main panels:

- R script:** The top-left pane shows the source code for a script named "01-Introduction.Rnw". The code includes comments and R commands for loading data, creating a plot, and calculating biomass. A text box labeled "R script" is overlaid on this pane.
- R console:** The bottom-left pane shows the execution of the R script. It displays the command prompt and the output of the commands, including the creation of the "kalimantan" data frame and the calculation of biomass. A text box labeled "R console" is overlaid on this pane.
- Graphical output:** The right-hand pane shows the "Environment" and "History" tabs. The "Environment" tab lists the objects in the global environment, including "hil.trees", "kal.plot", "kalimantan", "lsp.plots", "lsp", "pub", "web", and "valu". The "History" tab shows the execution history. Below these tabs, a box plot titled "Biomass estimation per plot with different models" is displayed. The y-axis is labeled "Biomass (Mg/ha)" and ranges from 100 to 500. The x-axis shows the different models: "Brown", "Yamakura", "Basuki", "Samalca", "Hashimoto", "Kenzo", "Ford", and "Jaya". A text box labeled "Graphical output" is overlaid on this pane.

- ▶ R Script/Viewer: Commands to execute, comments, data
- ▶ Console: Command line, type in command and hit Return
- ▶ Environment: All the dataframes, vector, variables, libraries and functions in your Workspace
- ▶ Output
 - ▶ Files in current directory
 - ▶ Plots (to export)
 - ▶ List of packages
 - ▶ Help files
 - ▶ Viewer to see data

- ▶ Designed for statistics
- ▶ 6000+ packages
- ▶ Freeware and Open Source
- ▶ Integrates with other programming languages like C, C++
- ▶ Constantly updated as new techniques emerge

You can download *R* and *R* Studio for free! You need both.

1. Download and install *R* (<https://cran.r-project.org/>) and *R* Studio (<https://www.rstudio.com/>)
2. Open *R* Studio
3. Do a tutorial series (<http://ditraglia.com/Econ103Public/>)

It is always a good idea to clear the environment before starting

```
rm(list=ls())
```

R needs to know where to look for things so we need to set the directory

```
getwd()
```

```
setwd("FolderPathGoesHere")
```


What are packages? Where to find them?

- ▶ Packages are bundles of commands, functions and data that you can use to do analysis
- ▶ R comes with a bunch of preloaded packages, but you may need more for the problem sets

Notable Examples:

- ▶ ggplot2: Famous package for making beautiful graphics
- ▶ quantmod: Tools for downloading financial data, plotting charts, and doing analysis
- ▶ stargazer: Package for creating beautiful regression tables

⇒ <https://cran.r-project.org>

What are packages? Where to find them?

- ▶ Packages are bundles of commands, functions and data that you can use to do analysis
- ▶ R comes with a bunch of preloaded packages, but you may need more for the problem sets

Notable Examples:

- ▶ `ggplot2`: Famous package for making beautiful graphics
- ▶ `quantmod`: Tools for downloading financial data, plotting charts, and doing analysis
- ▶ `stargazer`: Package for creating beautiful regression tables

⇒ <https://cran.r-project.org>

You can easily install them with *R* Studio, but you must load the package library each session so R knows you want to use it

Instructions:

1. Enter `install.packages("NameGoesHere")` into the Console.
2. Press the Return key
3. Add `library("NameGoesHere")` to the top of your R Script and run this command before starting each session

- ▶ An R Script is a text file (with the file extension `.R`) containing your R Code
- ▶ R Code is a series of comments, commands and functions that execute the tasks you require
- ▶ We can add comments using the `#` symbol

Where to find them?

You make your own :-)

- ▶ An R Script is a text file (with the file extension `.R`) containing your R Code
- ▶ R Code is a series of comments, commands and functions that execute the tasks you require
- ▶ We can add comments using the `#` symbol

Where to find them?

You make your own :-)

You can layer R Script together by telling R to use functions from another R Script.

Just add `Source("FileNameGoesHere.R")` at the top of your R Script and run it each session.

```
#####  
#  
# Econ 104 R Demo  
#  
#####  
  
# Set Directory and Install packages here  
setwd("/Users/korydkantenga/Dropbox/Teaching/Econometrics/")  
  
# Install  
#install.packages("maps")  
  
# Load Package  
library("maps")  
  
# Getting Help  
?mean    # help file  
??mean   # search
```

- ▶ Retype into Console and hit Return key line by line :-)
- ▶ Copy/Paste into Console and hit Return key line by line :-/
- ▶ Highlight a line and click Run :-)
- ▶ Click Source to run the entire script ;-)

More than likely you will run into errors. These errors are usually bugs.

Bugs are problems in **YOUR CODE**.
Almost all errors will be syntax errors.

Syntax errors are the coding equivalent of misspellings.

The 5 Stages of Debugging

At some point in each of our lives, we must face errors in our code. Debugging is a natural healing process to help us through these times. It is important to recognize these common stages and realize that debugging will eventually come to an end.



Denial

This stage is often characterized by such phrases as "What? That's impossible," or "I know this is right." A strong sign of denial is recompiling without changing any code, "just in case."



Bargaining/Self-Blame

Several programming errors are uncovered and the programmer feels stupid and guilty for having made them. Bargaining is common: "If I fix this, will you please compile?" Also, "I only have 14 errors to go!"



Anger

Cryptic error messages send the programmer into a rage. This stage is accompanied by an hours-long and profanity-filled diatribe about the limitations of the language directed at whomever will listen.



Depression

Following the outburst, the programmer becomes aware that hours have gone by unproductively and there is still no solution in sight. The programmer becomes listless. Posture often deteriorates.



Acceptance

The programmer finally accepts the situation, declares the bug a "feature", and goes to play some Quake.

1. Use `?x` to see the help file for a command with name `x`
2. Use `??x` to search for a command or package with name `x`
3. Enter `help.start()` into the Console and hit Return key
4. Try a tutorial series (<http://ditraglia.com/Econ103Public/>)
5. GIFY (G-It-For-Yourself)
6. Go to TA office hours
7. E-mail a TA

1. What is R?
2. Computing in R
3. Graphics in R
4. More in R

- ▶ Use `<-` and `=` to assign values (e.g. `x<-3.141593`)
- ▶ Standard arithmetic operators (e.g. `+`, `-`, `*`, `/`)
- ▶ Standard operators like
 - ▶ natural logarithm: `log(x)`
 - ▶ e^x : `exp(x)`
 - ▶ a^x : `a^x`

```
# Suppose we want to create a list of numbers.  
# We use c() for combine.  
x <- c(3,1,4,1,5,9,3)  
y <- c(2,7,1,8,2,8,2)  
  
# To create a column vector we combine rows.  
x <- rbind(3,1,4,1,5,9,3)  
print(x)
```

```
##      [,1]  
## [1,]    3  
## [2,]    1  
## [3,]    4  
## [4,]    1  
## [5,]    5  
## [6,]    9  
## [7,]    3
```

```
# To create a row vector we combine columns  
y <- cbind(2,7,1,8,2,8,2)  
print(y)
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7]  
## [1,]    2    7    1    8    2    8    2
```

```
# Transpose y into a column vector
y <- t(y)
# Combine in Data
data.new <- cbind(x,y)
# Turn into Data Frame
data.new <- as.data.frame(data.new)
print(head(data.new))
```

```
##   V1 V2
## 1  3  2
## 2  1  7
## 3  4  1
## 4  1  8
## 5  5  2
## 6  9  8
```

```
# Extract new dataset from Data Frame
data.sub <- subset(data.new, select = c("V2"))
print(head(data.sub))
```

```
##   V2
## 1  2
## 2  7
## 3  1
## 4  8
## 5  2
## 6  8
```

```
# Importing Data
caschool <- read.csv("caschool.csv")
print(head(caschool))
```

```
##      Observation.Number dist_cod  county                district
## 1                1      75119 Alameda          Sunol Glen Unified
## 2                2      61499  Butte          Manzanita Elementary
## 3                3      61549  Butte    Thermalito Union Elementary
## 4                4      61457  Butte Golden Feather Union Elementary
## 5                5      61523  Butte    Palermo Union Elementary
## 6                6      62042 Fresno    Burrel Union Elementary
##      gr_span enr1_tot teachers calw_pct meal_pct computer testscr comp_stu
## 1  KK-08      195    10.90   0.5102   2.0408        67  690.80 0.3435898
## 2  KK-08      240    11.15  15.4167  47.9167       101  661.20 0.4208333
## 3  KK-08     1550    82.90  55.0323  76.3226       169  643.60 0.1090323
## 4  KK-08      243    14.00  36.4754  77.0492        85  647.70 0.3497942
## 5  KK-08     1335    71.50  33.1086  78.4270       171  640.85 0.1280899
## 6  KK-08      137     6.40  12.3188  86.9565        25  605.55 0.1824818
##      expn_stu      str      avginc      el_pct read_scr math_scr
## 1 6384.911 17.88991 22.690001 0.000000   691.6   690.0
```



```
# List Variables
```

```
ls(caschool)
```

```
## [1] "avginc"           "calw_pct"         "comp_stu"
## [4] "computer"        "county"           "dist_cod"
## [7] "district"        "el_pct"           "enrl_tot"
## [10] "expn_stu"         "gr_span"          "math_scr"
## [13] "meal_pct"        "Observation.Number" "read_scr"
## [16] "str"             "teachers"         "testscr"
```

```
# Take logs of Average Income and
```

```
# add to data frame
```

```
caschool$logavginc <- log(caschool$avginc)
```

```
# Summary Statistics
```

```
mean(caschool$logavginc) #mean
```

```
## [1] 2.644841
```

```
sd(caschool$logavginc) #standard deviation
```

```
## [1] 0.392373
```

```
summary(caschool$logavginc) #quartiles
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      1.674   2.365   2.619   2.645   2.870   4.013
```

```
# Hypothesis Test on Mean  
# H0: mu = 2.6  
# H1: mu > 2.6  
xbar = mean(caschool$logavginc)  
serr = sd(caschool$logavginc)  
N     = length(caschool$logavginc)  
# test statistic (xbar-mu)/(sigma/sqrt(N))  
t = (xbar-2.6)/(serr/sqrt(N))  
# compute p-value from quantile function  
print(round(1-pnorm(t),2))
```


```
## [1] 0.01
```

```
# Hence, we fail to reject  
# the null at 0.5% significance level  
# but may reject at the 1% level.
```


1. What is R?
2. Computing in R
3. Graphics in R
4. More in R

```
# Is log average income normal?  
hist(caschool$logavginc,xlab="",main="Log of Average Income")
```

Log of Average Income



```
# Is there a relationship between income and test scores?  
plot(caschool$avginc,caschool$testscr, xlab="Average Income ($)",ylab="Test Score"  
 , main="Income v. Scores")
```




```
# Does income predict test scores?
# Simple Linear Regression
lm.ols <- lm(testscr ~ avginc, data = caschool)
summary(lm.ols)

##
## Call:
## lm(formula = testscr ~ avginc, data = caschool)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -39.574  -8.803   0.603   9.032  32.530
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  625.3836    1.5324   408.11  <2e-16 ***
## avginc        1.8785     0.0905   20.76  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 13.39 on 418 degrees of freedom
## Multiple R-squared:  0.5076, Adjusted R-squared:  0.5064
## F-statistic: 430.8 on 1 and 418 DF,  p-value: < 2.2e-16
```

```
plot(caschool$avginc,caschool$testscr, xlab="Average Income ($)",ylab="Test Score"  
 , main="Income v. Scores")  
abline(lm(testscr ~ avginc, data = caschool),col="red")
```

Income v. Scores



1. What is R?
2. Computing in R
3. Graphics in R
4. More in R

- ▶ Functions
- ▶ Loops
- ▶ Low-level Coding
- ▶ Astral Projection and Telekinesis

For all of these and more, try some *R* tutorials and GIFY.

1. Descriptive Statistics (e.g. mean, median, mode)
2. Probability Theory (e.g. Bayes Rule)
3. Random Variables (Discrete, Continuous)
4. Probability Distributions (e.g. CDF, PDF, PMF)
5. Expectations (e.g. $\mathbb{E}[aX] = a \cdot \mathbb{E}[X]$)
6. Point and Interval Estimation (e.g. confidence intervals)
7. Properties (e.g. unbiasedness, efficiency, consistency)
8. Sampling Distributions (e.g. $\bar{X} \sim N(\mu, \sigma^2/N)$)
9. Hypothesis Testing (e.g. $H_0 : \mu = 0, H_1 : \mu \neq 0$)
10. Simple Linear Regression (e.g. $y_i = \alpha + \beta x_i + \varepsilon_i$)

1. Descriptive Statistics (e.g. mean, median, mode)
2. Probability Theory (e.g. Bayes Rule)
3. Random Variables (Discrete, Continuous)
4. Probability Distributions (e.g. CDF, PDF, PMF)
5. Expectations (e.g. $\mathbb{E}[aX] = a \cdot \mathbb{E}[X]$)
6. Point and Interval Estimation (e.g. confidence intervals)
7. Properties (e.g. unbiasedness, efficiency, consistency)
8. Sampling Distributions (e.g. $\bar{X} \sim N(\mu, \sigma^2/N)$)
9. Hypothesis Testing (e.g. $H_0 : \mu = 0, H_1 : \mu \neq 0$)
10. Simple Linear Regression (e.g. $y_i = \alpha + \beta x_i + \varepsilon_i$)

1. Descriptive Statistics (e.g. mean, median, mode)
2. Probability Theory (e.g. Bayes Rule)
- 3. Random Variables (Discrete, Continuous)**
4. Probability Distributions (e.g. CDF, PDF, PMF)
5. Expectations (e.g. $\mathbb{E}[aX] = a \cdot \mathbb{E}[X]$)
6. Point and Interval Estimation (e.g. confidence intervals)
7. Properties (e.g. unbiasedness, efficiency, consistency)
8. Sampling Distributions (e.g. $\bar{X} \sim N(\mu, \sigma^2/N)$)
9. Hypothesis Testing (e.g. $H_0 : \mu = 0, H_1 : \mu \neq 0$)
10. Simple Linear Regression (e.g. $y_i = \alpha + \beta x_i + \varepsilon_i$)

1. Descriptive Statistics (e.g. mean, median, mode)
2. Probability Theory (e.g. Bayes Rule)
3. Random Variables (Discrete, Continuous)
- 4. Probability Distributions (e.g. CDF, PDF, PMF)**
5. Expectations (e.g. $\mathbb{E}[aX] = a \cdot \mathbb{E}[X]$)
6. Point and Interval Estimation (e.g. confidence intervals)
7. Properties (e.g. unbiasedness, efficiency, consistency)
8. Sampling Distributions (e.g. $\bar{X} \sim N(\mu, \sigma^2/N)$)
9. Hypothesis Testing (e.g. $H_0 : \mu = 0, H_1 : \mu \neq 0$)
10. Simple Linear Regression (e.g. $y_i = \alpha + \beta x_i + \varepsilon_i$)

1. Descriptive Statistics (e.g. mean, median, mode)
2. Probability Theory (e.g. Bayes Rule)
3. Random Variables (Discrete, Continuous)
4. Probability Distributions (e.g. CDF, PDF, PMF)
- 5. Expectations (e.g. $\mathbb{E}[aX] = a \cdot \mathbb{E}[X]$)**
6. Point and Interval Estimation (e.g. confidence intervals)
7. Properties (e.g. unbiasedness, efficiency, consistency)
8. Sampling Distributions (e.g. $\bar{X} \sim N(\mu, \sigma^2/N)$)
9. Hypothesis Testing (e.g. $H_0 : \mu = 0, H_1 : \mu \neq 0$)
10. Simple Linear Regression (e.g. $y_i = \alpha + \beta x_i + \varepsilon_i$)

1. Descriptive Statistics (e.g. mean, median, mode)
2. Probability Theory (e.g. Bayes Rule)
3. Random Variables (Discrete, Continuous)
4. Probability Distributions (e.g. CDF, PDF, PMF)
5. Expectations (e.g. $\mathbb{E}[aX] = a \cdot \mathbb{E}[X]$)
- 6. Point and Interval Estimation (e.g. confidence intervals)**
7. Properties (e.g. unbiasedness, efficiency, consistency)
8. Sampling Distributions (e.g. $\bar{X} \sim N(\mu, \sigma^2/N)$)
9. Hypothesis Testing (e.g. $H_0 : \mu = 0, H_1 : \mu \neq 0$)
10. Simple Linear Regression (e.g. $y_i = \alpha + \beta x_i + \varepsilon_i$)

1. Descriptive Statistics (e.g. mean, median, mode)
2. Probability Theory (e.g. Bayes Rule)
3. Random Variables (Discrete, Continuous)
4. Probability Distributions (e.g. CDF, PDF, PMF)
5. Expectations (e.g. $\mathbb{E}[aX] = a \cdot \mathbb{E}[X]$)
6. Point and Interval Estimation (e.g. confidence intervals)
- 7. Properties (e.g. unbiasedness, efficiency, consistency)**
8. Sampling Distributions (e.g. $\bar{X} \sim N(\mu, \sigma^2/N)$)
9. Hypothesis Testing (e.g. $H_0 : \mu = 0, H_1 : \mu \neq 0$)
10. Simple Linear Regression (e.g. $y_i = \alpha + \beta x_i + \varepsilon_i$)

1. Descriptive Statistics (e.g. mean, median, mode)
2. Probability Theory (e.g. Bayes Rule)
3. Random Variables (Discrete, Continuous)
4. Probability Distributions (e.g. CDF, PDF, PMF)
5. Expectations (e.g. $\mathbb{E}[aX] = a \cdot \mathbb{E}[X]$)
6. Point and Interval Estimation (e.g. confidence intervals)
7. Properties (e.g. unbiasedness, efficiency, consistency)
- 8. Sampling Distributions (e.g. $\bar{X} \sim N(\mu, \sigma^2/N)$)**
9. Hypothesis Testing (e.g. $H_0 : \mu = 0, H_1 : \mu \neq 0$)
10. Simple Linear Regression (e.g. $y_i = \alpha + \beta x_i + \varepsilon_i$)

1. Descriptive Statistics (e.g. mean, median, mode)
2. Probability Theory (e.g. Bayes Rule)
3. Random Variables (Discrete, Continuous)
4. Probability Distributions (e.g. CDF, PDF, PMF)
5. Expectations (e.g. $\mathbb{E}[aX] = a \cdot \mathbb{E}[X]$)
6. Point and Interval Estimation (e.g. confidence intervals)
7. Properties (e.g. unbiasedness, efficiency, consistency)
8. Sampling Distributions (e.g. $\bar{X} \sim N(\mu, \sigma^2/N)$)
- 9. Hypothesis Testing (e.g. $H_0 : \mu = 0, H_1 : \mu \neq 0$)**
10. Simple Linear Regression (e.g. $y_i = \alpha + \beta x_i + \varepsilon_i$)

1. Descriptive Statistics (e.g. mean, median, mode)
2. Probability Theory (e.g. Bayes Rule)
3. Random Variables (Discrete, Continuous)
4. Probability Distributions (e.g. CDF, PDF, PMF)
5. Expectations (e.g. $\mathbb{E}[aX] = a \cdot \mathbb{E}[X]$)
6. Point and Interval Estimation (e.g. confidence intervals)
7. Properties (e.g. unbiasedness, efficiency, consistency)
8. Sampling Distributions (e.g. $\bar{X} \sim N(\mu, \sigma^2/N)$)
9. Hypothesis Testing (e.g. $H_0 : \mu = 0, H_1 : \mu \neq 0$)
10. Simple Linear Regression (e.g. $y_i = \alpha + \beta x_i + \varepsilon_i$)

1. Linear Regression Models

- ▶ Least Squares Estimation
- ▶ Inference

2. Panel Data Models

- ▶ First Difference
- ▶ Fixed Effects

3. Simultaneous-Equations Models

- ▶ Instrumental Variables

4. Discrete Choice Models

- ▶ Maximum Likelihood Estimation

5. Time Series Models

- ▶ Xie, Yihui (2014). *An Introduction to R*. UPenn, The Warren Center. Nov 21. <http://ditraglia.com/pdf/yihui1.pdf>
- ▶ The R Project for Statistical Computing.
<https://www.r-project.org>
- ▶ R-Bloggers. <https://www.r-bloggers.com>